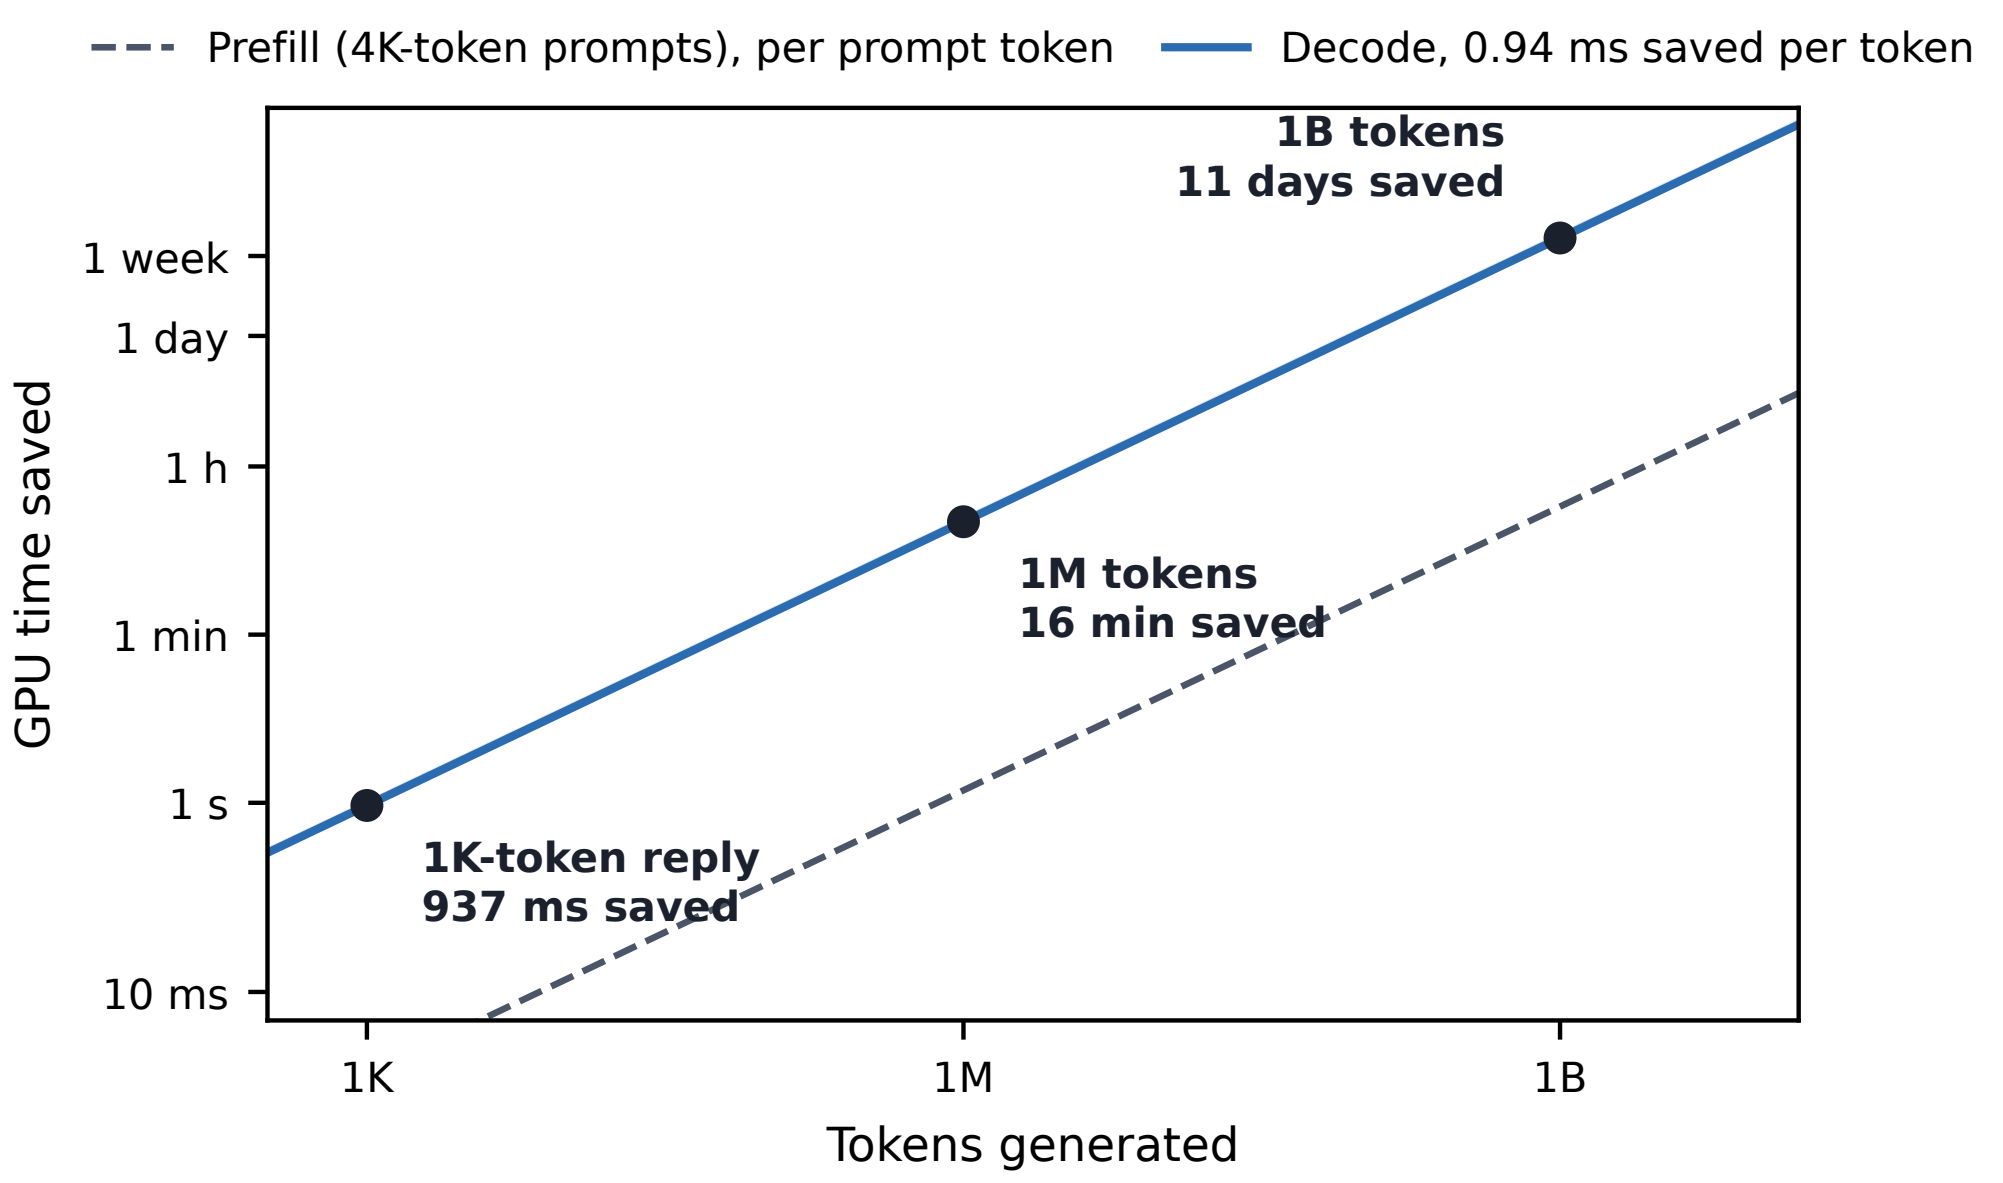
Prefill (4K-token prompts), per prompt token
Decode, 0.94 ms saved per token
1B tokens
11 days saved
1M tokens
16 min saved
1K-token reply
937 ms saved
1 week
1 day
1 h
1 min
1 s
10 ms
GPU time saved
1K
1M
1B
Tokens generated